

DEEPAK KUMAR SARANGI

Senior Software Engineer | Distributed Systems & Data Platforms

Bangalore | 8270010477 | dpksarangi@gmail.com | [linkedin.com/in/deepak-sarangi](https://www.linkedin.com/in/deepak-sarangi) | portfolio-v21z.vercel.app

PROFESSIONAL SUMMARY

Senior Software Engineer with 10+ years of experience architecting high-throughput, event-driven data platforms at scale. At Walmart, designed and owned pipelines processing billions of events daily on GCP, cutting data delivery latency by 90%. Strong track record of driving cross-team technical decisions, authoring RFCs adopted org-wide, and building GenAI-powered internal tools. Targeting Staff / Lead Engineer roles where I can own system design, define engineering standards, and lead platform-level initiatives.

EXPERIENCE

Walmart | Senior Software Engineer

Feb 2022 – Present

Platform Architecture & Data Engineering

- Designed and owned end-to-end architecture for event-driven data pipelines processing 1–10 billion events/day on GCP (Pub/Sub, Dataflow, BigQuery, Bigtable) — defined SLAs, failure handling, and scalability patterns adopted across 3+ downstream teams.
- Architected a Kafka + Spark streaming ingestion platform that reduced data delivery latency from 3 hours to 30 minutes (90% improvement), directly unblocking downstream ML and analytics consumers.
- Built a Spring Boot–based multi-tenant data distribution platform integrating GCS, Hive, and SFTP/API clients with resiliency patterns (circuit breakers, retries, dead-letter queues) for fault-tolerant large-scale data transfer.
- Owned 5–10 production data pipelines end-to-end; optimized Spark job runtime by 20% through query tuning, partitioning strategy, and execution plan analysis.
- Implemented data validation and quality enforcement systems, reducing data defects by 25% across pipelines.

GenAI & LLM Tooling

- Built an end-to-end LLM-powered developer productivity agent that converts Jira requirements into production-ready PRs: scans requirements → plans → generates code → builds → auto-fixes CI failures → deploys → runs tests → raises a merge-ready PR with zero manual handoff.
- Designed an LLM-powered cost intelligence tool for internal teams: scans AI/LLM project codebases, detects cost-leakage patterns (context rot, suboptimal prompt compression, oversized RAG fetch windows, high-cost model overuse), quantifies potential savings, validates fixes, and raises merge-ready PRs in target codebases.
- Built multi-tenant chatbot architecture with tenant isolation and shared services, reducing infrastructure cost by 40% while improving self-service operations.

Technical Leadership & Cross-Team Influence

- Authored RFCs and system design documents for 3+ cross-team initiatives; standards adopted across CRM and data platform organizations.
- Drove cross-team technical decisions on streaming architecture, data modeling, and pipeline reliability — key voice in design reviews across squads.
- Mentored 3–5 engineers across experience levels; led hiring panels and technical interviews.
- Held end-to-end ownership over 2 of 5 core applications within the team; primary design decision partner for the team lead, contributing to architecture and technical direction across all 5 apps.
- Presented internal tech forums on streaming platform design and LLM-powered developer tooling.

Fidelity Investments | Lead Software Engineer

May 2021 – Feb 2022

- Contributed to design and development of a pricing engine supporting batch and near-real-time data processing for financial instruments.
- Integrated real-time data feeds from multiple upstream sources, improving data freshness and system responsiveness.
- Implemented observability stack with metrics, dashboards, and log analysis (Prometheus/Grafana), improving incident detection and response time.

Tata Consultancy Services (TCS) | IT Analyst

Mar 2016 – May 2021

Progression: Assistant System Engineer → System Engineer → IT Analyst | 3 client projects across multiple domains

- Designed and built Java backend services from scratch across 3 client projects spanning multiple domains, taking ownership of requirements through to deployment on each engagement.

- Developed and maintained enterprise middleware and integration layers, enabling reliable communication across heterogeneous systems in large-scale client environments.
- Identified and resolved production incidents independently through structured root cause analysis, improving system stability and reducing recurrence across integration workflows.
- Drove performance improvements in backend integration workflows, reducing processing latency and improving throughput for downstream consumers.
- Mentored and onboarded junior engineers joining the team, providing technical guidance on Java development, integration patterns, and debugging practices.
- Progressed through three internal roles over 5 years — a recognition of consistent delivery, cross-project adaptability, and growing ownership across client engagements.

TECHNICAL SKILLS

Languages: Java, Python, JavaScript, SQL

Big Data & Streaming: Apache Spark, Apache Kafka, Apache Airflow, Pub/Sub, Dataflow

Cloud (GCP): BigQuery, Bigtable, Cloud Storage, Dataflow, Pub/Sub, Dataproc, GCS

Databases: PostgreSQL, MySQL, Cassandra, Cosmos DB

Backend & APIs: Spring Boot, RESTful APIs, Microservices, Event-Driven Architecture

GenAI & LLM: LangChain, OpenAI APIs, Claude, Ollama, Prompt Engineering, RAG

Observability: Prometheus, Grafana, Splunk, Openobserve

DevOps & CI/CD: Docker, Kubernetes, Jenkins, Concord, Git

EDUCATION

Bachelor of Technology, Mechanical Engineering | GCEK, Bhawanipatna | CGPA: 7.7